

Transaction Management

- **Transaction Concept**
- **A Simple Transaction Model**
- **Serializability**

Transaction Concept

- Often, a collection of several operations on the database appears to be a single unit from the point of view of the database user.
- For example,
“a transfer of funds from a checking account to a savings account”
- is a single operation from the customer’s standpoint; within the database system,
- however, it consists of several operations.
- Clearly, it is essential that all these operations occur, or that, in case of a failure, none occur.
- Collections of operations that form a single logical unit of work are called **transactions**.

Transaction Concept

- A **transaction** is a **unit** of program execution that accesses and possibly updates various data items.
- Usually, a transaction is initiated by a user program written in a high-level data-manipulation language (typically SQL), or programming language (for example, C++, or Java), with embedded database accesses in JDBC or ODBC.
- A transaction is delimited by statements (or function calls) of the form **begin transaction** and **end transaction**. The transaction consists of all operations executed between the **begin transaction** and **end transaction**.

ACID Property

- **Atomicity.** Either all operations of the transaction are reflected properly in the database, or none are.
- **Consistency.** Execution of a transaction in isolation (that is, with no other transaction executing concurrently) preserves the consistency of the database.
- **Isolation.** Even though multiple transactions may execute concurrently, the system guarantees that, for every pair of transactions T_i and T_j , it appears to T_i that either T_j finished execution before T_i started or T_j started execution after T_i finished. Thus, each transaction is unaware of other transactions executing concurrently in the system.
- **Durability.** After a transaction completes successfully, the changes it has made to the database persist, even if there are system failures.

A Simple Transaction Model

- The data items in our simplified model contain a single data value (a number in our examples).
- Each data item is identified by a name (typically a single letter in our examples, **that is, A, B, C, etc.**).
- We shall illustrate the transaction concept using a simple bank application consisting of several accounts and a set of transactions that access and update those accounts.
- Transactions access data using two operations:
 - **read(X)** - which transfers the data item X from the database to a variable, also called X, in a buffer in main memory belonging to the transaction that executed the read operation.
Eg. read(Empno) means, the variable Empno = value in the attribute Empno.
 - **write(X)**, which transfers the value in the variable X in the main-memory buffer of the transaction that executed the write to the data item X in the database.
Eg. write(Empno) means, the attribute Empno = value in the variable Empno

A Simple Transaction Model

- Let T_i be a transaction that transfers \$50 from account A to account B . This transaction can be defined as:

Consistency: The consistency requirement here is that the sum of A and B be unchanged by the execution of the transaction.

Atomicity: Suppose $A = 1000\$$ and $B = 2000\$$.

If there is an error after the $\text{write}(A)$ operation but before the $\text{write}(B)$ then $A = 950\$$ and $B = 2000\$$.

the sum $A + B$ is no longer preserved.

Now the system is in inconsistent state.

Durability: Once the execution of the transaction completes successfully, it must be the case that no system failure can result in a loss of data corresponding to this transfer of funds.

Isolation: if several transactions are executed concurrently, their operations may interleave in some undesirable way, resulting in an inconsistent state.

Eg. At time when $A-50$ is done but not $B+50$, if an other transaction reads A and B .

Ensuring atomicity & Durability is the responsibility of the database system; specifically, it is handled by a component of the database called the **recovery system**.

Ensuring the isolation : by **concurrency-control system**.


T_i : $\text{read}(A);$
 $A := A - 50;$
 $\text{write}(A);$
 $\text{read}(B);$
 $B := B + 50;$
 $\text{write}(B).$

A Simple Transaction Model

- a transaction may not always complete its execution successfully, Such a transaction is termed **aborted**.
- Once the changes caused by an aborted transaction have been undone, we say that the transaction has been **rolled back**.
- It is part of the responsibility of the recovery scheme to manage transaction aborts. This is done typically by maintaining a **log**.
- A transaction that completes its execution successfully is said to be **committed**.
- Once a transaction has committed, we cannot undo its effects by aborting it. The only way to undo the effects of a committed transaction is to execute **a compensating transaction**.
 - For instance, if a transaction added \$20 to an account, the compensating transaction would subtract \$20 from the account.
- However, it is not always possible to create such a compensating transaction. Therefore, the responsibility of writing and executing a compensating transaction is left to the user, and is not handled by the database system.

A Simple Abstract Transaction Model

- **Active**, the initial state; the transaction stays in this state while it is executing.
- **Partially committed**, after the final statement has been executed.
- **Failed**, after the discovery that normal execution can no longer proceed.
- **Aborted**, after the transaction has been rolled back and the database has been restored to its state prior to the start of the transaction.
- **Committed**, after successful completion.




State diagram of a transaction.

A transaction starts in the **active state**.

When it finishes its final statement, it enters the **partially committed state**. At this point, the transaction has completed its execution, but it is still possible that it may have to be aborted, since the actual output may still be temporarily residing in main memory, and thus a hardware failure may preclude its successful completion.

A transaction is said to have **terminated** if it has either committed or aborted.

A Simple Transaction Model



A transaction enters the **failed** state after the system determines that the transaction can no longer proceed with its normal execution (for example, because of hardware or logical errors). Such a transaction must be **rolled back**. Then, it enters the aborted state.

At this point, the system has two options:

- It can **restart** the transaction, but only if the transaction was aborted as a result of some hardware or software error that was not created through the internal logic of the transaction. A restarted transaction is considered to be a new transaction.
- It can **kill** the transaction. It usually does so because of some internal logical error that can be corrected only by rewriting the application program, or because the input was bad, or because the desired data were not found in the database.

Transaction Isolation

- Transaction-processing systems usually allow **multiple transactions to run concurrently**.
- Allowing multiple transactions to update data concurrently causes several complications with consistency of the data, as we saw earlier.
- Ensuring consistency in spite of concurrent execution of transactions requires extra work;
- It is far easier to insist that transactions run **serially**—that is, one at a time, each starting only after the previous one has completed.
- However, there are two good reasons for allowing concurrency:
 - (1) **Improved throughput and resource utilization.**
 - (2) **Reduced waiting time.**
- The database system must control the interaction among the concurrent transactions to prevent them from destroying the consistency of the database. It does so through a variety of mechanisms called **concurrency-control scheme**.

A Simple Transaction Model

Consider the simplified banking system which has several accounts, and a set of transactions that access and update those accounts.

Let T_1 and T_2 be two transactions that transfer funds from one account to another.

Transaction T_1 transfers \$50 from account A to account B .

Transaction T_2 transfers 10 percent of the balance from account A to account B .

```
 $T_1$ : read( $A$ );  
       $A := A - 50$ ;  
      write( $A$ );  
      read( $B$ );  
       $B := B + 50$ ;  
      write( $B$ ).
```

```
 $T_2$ : read( $A$ );  
       $temp := A * 0.1$ ;  
       $A := A - temp$ ;  
      write( $A$ );  
      read( $B$ );  
       $B := B + temp$ ;  
      write( $B$ ).
```

A Simple Transaction Model

- Suppose the current values of accounts *A* and *B* are \$1000 and \$2000, respectively.
- Suppose also that the two transactions are executed one at a time in the order *T*₁ followed by *T*₂. (**Schedule #1**)

<i>T</i> ₁	<i>T</i> ₂
read(<i>A</i>) <i>A</i> := <i>A</i> − 50 write(<i>A</i>) read(<i>B</i>) <i>B</i> := <i>B</i> + 50 write(<i>B</i>) commit	read(<i>A</i>) <i>temp</i> := <i>A</i> * 0.1 <i>temp</i> = 95 <i>A</i> := <i>A</i> − <i>temp</i> write(<i>A</i>) read(<i>B</i>) <i>B</i> := <i>B</i> + <i>temp</i> write(<i>B</i>) commit
<i>A</i> = 950, <i>B</i> = 2050	<i>A</i> = 855, <i>B</i> = 2145
	<i>A</i> + <i>B</i> = 3000

Schedule 1 — a serial schedule in which *T*₁ is followed by *T*₂.

The final values of accounts *A* and *B*, after the execution of this schedule are \$855 and \$2145, respectively.

Thus, the total amount of money in accounts *A* and *B*—that is, the **sum *A* + *B***—is **preserved** after the execution of both transactions.

A Simple Transaction Model

- Similarly, if the transactions are executed one at a time in the order T_2 followed by T_1 , the sum $A + B$ is preserved, and the final values of accounts A and B are \$850 and \$2150, respectively.

(Schedule #2)

A and B are \$1000 and \$2000,

T_1	T_2
read(A) $A := A - 50$ write(A) read(B) $B := B + 50$ write(B) commit	read(A) $temp := A * 0.1$ $temp = 100$ $A := A - temp$ write(A) read(B) $B := B + temp$ write(B) commit

$A = 900,$
 $B = 2100$

$A = 850,$
 $B = 2150$

$A+B = 3000$

Schedule 2 — a serial schedule in which T_2 is followed by T_1 .

- The execution sequences just described are called **schedules**.
- They represent the chronological order in which instructions are executed in the system.
- Clearly, a schedule for a set of transactions must consist of all instructions of those transactions, and must preserve the order in which the instructions appear in each individual transaction.
For example, in transaction T_1 , the instruction **write(A)** must appear before the instruction **read(B)**, in any valid schedule.
- These schedules are **serial**.
- Each serial schedule consists of a sequence of instructions from various transactions, where the instructions belonging to one single transaction appear together in that schedule.
- for a set of n transactions, there exist n factorial ($n!$) different valid **serial schedules**.

Concurrent Execution of Transactions

- When the database system executes **several transactions** concurrently, the corresponding schedule no longer needs to be serial.
- If two transactions are running concurrently, the operating system may execute one transaction for a little while, then perform a **context switch**, execute the second transaction for some time, and then switch back to the first transaction for some time, and so on.
- With multiple transactions, the **CPU time is shared** among all the transactions.
- Several execution sequences are possible, since the various instructions from both transactions may now be **interleaved**.
- In general, it is not possible to predict exactly how many instructions of a transaction will be executed before the CPU switches to another transaction.

- Suppose that the two transactions (T1 and T2) are executed concurrently, One possible schedule is

A and B are \$1000 and \$2000,

T_1	T_2
read(A) $A := A - 50$ write(A) read(B) $B := B + 50$ write(B) commit	(Schedule #1) read(A) $temp := A * 0.1$ temp = 95 $A := A - temp$ write(A) read(B) $B := B + temp$ write(B) commit

A = 950,
B = 2050

A = 855,
B = 2145

Schedule 1 — a serial schedule in which T_1 is followed by T_2 .

A+B = 3000

(Schedule #3)

T_1	T_2
read(A) $A := A - 50$ write(A) read(B) $B := B + 50$ write(B) commit	read(A) $temp := A * 0.1$ temp = 95 $A := A - temp$ write(A) read(B) $B := B + temp$ write(B) commit

A = 950

B = 2050

A = 855

B = 2145

A = 855

A+B = 3000

Schedule 3 — a concurrent schedule equivalent to schedule 1.

After this execution takes place, we arrive at the same state as the one in which the transactions are executed serially in the order T_1 followed by T_2 .

The sum $A + B$ is indeed preserved.

- Not all concurrent executions result in a correct state. To illustrate, consider the schedule.

(Schedule #3)

A and B are \$1000 and \$2000

T ₁	T ₂
read(A) A := A − 50 write(A)	read(A) temp := A * 0.1 temp = 95 A := A − temp write(A) A = 855
read(B) B := B + 50 write(B) commit	read(B) B := B + temp write(B) B = 2145 commit A = 855

Schedule 3 — a concurrent schedule equivalent to schedule 1.

(Schedule #4)

T ₁	T ₂
read(A) Var.A = 950 A := A − 50	read(A) A = 1000 temp := A * 0.1 temp = 100 A := A − temp write(A) A = 900 read(B) Var. B = 2000
A = 950 write(A) read(B) B := B + 50 B = 2050 write(B) commit	B := B + temp write(B) B = 2100 commit A = 950

Schedule 4 — a concurrent schedule resulting in an inconsistent state.

A+B = 3050

A + B is not preserved
 This final state is an *inconsistent state*,

- If control of concurrent execution is left entirely to the **operating system**, many possible schedules, including ones that leave the database in an inconsistent state, are possible.
- It is the job of the **database system** to ensure that any schedule that is executed will leave the database in a consistent state.
- The **concurrency-control component** of the database system carries out this task.
- We can ensure consistency of the database under concurrent execution by making sure that any schedule that is executed has the same effect as a schedule that could have occurred without any concurrent execution.
 - That is, the schedule should, in some sense, be equivalent to a serial schedule.
 - Such schedules are called **serializable** schedules.

Serializability

- how to determine when a schedule is **serializable**?
 - Certainly, serial schedules are **serializable**,
 - but if steps of multiple transactions are interleaved, it is harder to determine whether a schedule is serializable.
- Since transactions are programs, it is difficult to determine exactly what operations a transaction performs and how operations of various transactions interact.
- To simplify, we consider only the **read** and **write** operations in the schedule.

T ₁	T ₂
read(A) A := A - 50 write(A)	read(A) temp := A * 0.1 A := A - temp write(A)
read(B) B := B + 50 write(B) commit	read(B) B := B + temp write(B) commit

Schedule 3 — a concurrent schedule equivalent to schedule 1.

T ₁	T ₂
read(A) write(A)	read(A) write(A)
read(B) write(B)	read(B) write(B)

Schedule 3 — showing only the read and write instructions.

Serializability

- There are different forms of schedule equivalence, but focus on a particular form called **conflict serializability**.
- Let us consider a schedule S in which there are two consecutive instructions, I and J , of transactions T_i and T_j , respectively ($i \neq j$).
- If I and J refer to **different data items**, then we can swap I and J without affecting the results of any instruction in the schedule.
- However, if I and J refer to the **same data item** Q , then the order of the two steps may matter.
- Since we are dealing with only read and write instructions, there are four cases that we need to consider:

Serializability

- Since we are dealing with only read and write instructions, there are four cases that we need to consider:

1. $I = \text{read}(Q), J = \text{read}(Q)$. The order of I and J does not matter, since the same value of Q is read by T_i and T_j , regardless of the order.
2. $I = \text{read}(Q), J = \text{write}(Q)$. If I comes before J , then T_i does not read the value of Q that is written by T_j in instruction J . If J comes before I , then T_i reads the value of Q that is written by T_j . Thus, the order of I and J matters.
3. $I = \text{write}(Q), J = \text{read}(Q)$. The order of I and J matters for reasons similar to those of the previous case.
4. $I = \text{write}(Q), J = \text{write}(Q)$. Since both instructions are write operations, the order of these instructions does not affect either T_i or T_j . However, the value obtained by the next $\text{read}(Q)$ instruction of S is affected, since the result of only the latter of the two write instructions is preserved in the database. If there is no other $\text{write}(Q)$ instruction after I and J in S , then the order of I and J directly affects the final value of Q in the database state that results from schedule S .

Thus, only in the case where both I and J are read instructions does the relative order of their execution not matter.

We say that I and J **conflict** if they are operations by different transactions on the same data item, and at least one of these instructions is a **write** operation.

Serializability

Illustration

schedule 3	
T_1	T_2
read(A) write(A)	read(A) write(A)
read(B) write(B)	

Schedule 5	
T_1	T_2
read(A) write(A)	read(A) write(A)
read(B)	
write(B)	
	read(B) write(B)

- The **write(A)** instruction of T_1 conflicts with the **read(A)** instruction of T_2 .
- However, the **write(A)** instruction of T_2 does not conflict with the **read(B)** instruction of T_1 , because the two instructions access different data items.
- Let I and J be consecutive instructions of a schedule S .
- If I and J are instructions of different transactions and I and J do not conflict, then we can swap the order of I and J to produce a new schedule $S1$.
- $S1$ is equivalent to S , since all instructions appear in the same order in both schedules except for I and J , whose order does not matter.

Since the **write(A)** instruction of T_2 in schedule 3 does not conflict with the **read(B)** instruction of T_1 , we can swap these instructions to generate an equivalent schedule, schedule 5.

Regardless of the initial system state, schedules 3 and 5 both produce the same final system state.

Serializability

schedule 3		Schedule 5						Schedule 6		
T_1	T_2	T_1	T_2	T1	T2		T1	T2	T_1	T_2
read(A)	read(A) write(A)	read(A)	read(A)	read(A)			read(A)		read(A)	read(A) write(A) read(B) write(B)
write(A)		write(A)		write(A)			write(A)			
read(B)		read(B)			read(A)			read(A)		
write(B)		write(B)			write(A)			write(B)		
	read(B)		read(B)	write(B)				write(A)		read(A)
	write(B)		write(B)		read(B)			read(B)		write(A)
					write(B)			write(B)		read(B)
										write(B)

- We continue to swap nonconflicting instructions:
 - Swap the read(B) instruction of T1 with the read(A) instruction of T2.
 - Swap the write(B) instruction of T1 with the write(A) instruction of T2.
 - Swap the write(B) instruction of T1 with the read(A) instruction of T2.

Schedule 6 is equivalent to Schedule 1.

So, Schedule 3 is serializable schedule.

Schedule 1	
T ₁	T ₂
read(A) A := A - 50 write(A) read(B) B := B + 50 write(B) commit	read(A) temp := A * 0.1 A := A - temp write(A) read(B) B := B + temp write(B) commit

Serializability

- If a schedule S can be transformed into a schedule S by a series of swaps of nonconflicting instructions, we say that S and S are **conflict equivalent**.
- Not all serial schedules are conflict equivalent to each other.
- For example, schedules 1 and 2 are not conflict equivalent.
- The concept of conflict equivalence leads to the concept of conflict serializability.
- We say that a schedule S is **conflict serializable** if it is conflict equivalent to a serial schedule.
- Thus, schedule 3 is conflict serializable, since it is conflict equivalent to the serial schedule 1.

Schedule 4 is not conflict serializable

T_1	T_2
read(A) $A := A - 50$	read(A) $temp := A * 0.1$ $A := A - temp$ write(A) read(B)
write(A) read(B) $B := B + 50$ write(B) commit	
	$B := B + temp$ write(B) commit

T1	T2	T1	T2	T1	T2	T1	T2
read(A)			read(A)		read(A)		read(A)
	read(A)	read(A)		read(A)		read(A)	
	write(A)		write(A)		write(A)		write(A)
	read(B)		read(B)	write(A)		write(A)	
write(A)		write(A)			read(B)	read(B)	
read(B)		read(B)		read(B)			read(B)
write(B)		write(B)		write(B)		write(B)	
	Write(B)		Write(B)		Write(B)		Write(B)

Serializability

- Finally, consider schedule 7, it consists of only the significant operations (that is, the read and write) of transactions T_3 and T_4 .
- This schedule is not conflict serializable, since it is not equivalent to either the serial schedule $\langle T_3, T_4 \rangle$ or the serial schedule $\langle T_4, T_3 \rangle$.

Schedule 7.

T_3	T_4
read(Q)	write(Q)
write(Q)	

Conflict Serializability – by precedence graph

- a simple and efficient method for determining conflict serializability of a schedule.
- Consider a schedule S .
- We construct a directed graph, called a **precedence graph**, from S . This graph consists of a pair $G = (V, E)$, where V is a set of vertices and E is a set of edges.
- The set of vertices consists of all the transactions participating in the schedule. The set of edges consists of all edges $T_i \rightarrow T_j$ for which one of three conditions holds:

T_1	T_2
read(A) write(A)	read(A) write(A)
read(B) write(B)	read(B) write(B)

- 1. T_i executes write(Q) before T_j executes read(Q).
- 2. T_i executes read(Q) before T_j executes write(Q).
- 3. T_i executes write(Q) before T_j executes write(Q).


If an edge $T_i \rightarrow T_j$ exists in the precedence graph, then, in any serial schedule S equivalent to S , T_i must appear before T_j .

Conflict Serializability – by precedence graph

Example

Schedule 1

T_1	T_2
read(A) $A := A - 50$ write(A) read(B) $B := B + 50$ write(B) commit	read(A) $temp := A * 0.1$ $A := A - temp$ write(A) read(B) $B := B + temp$ write(B) commit



Precedence graph for (a) schedule 1 and (b) schedule 2.

Schedule 2

T_1	T_2
read(A) $A := A - 50$ write(A) read(B) $B := B + 50$ write(B) commit	read(A) $temp := A * 0.1$ $A := A - temp$ write(A) read(B) $B := B + temp$ write(B) commit


For example, the precedence graph for schedule 1 is in (a) contains the single edge $T_1 \rightarrow T_2$, since all the instructions of T_1 are executed before the first instruction of T_2 is executed.

Similarly, (b) shows the precedence graph for schedule 2 with the single edge $T_2 \rightarrow T_1$, since all the instructions of T_2 are executed before the first instruction of T_1 is executed.

Schedule 3

T_1	T_2
read(A) write(A)	read(A) write(A)
read(B) write(B)	read(B) write(B)


Precedence graph of Schedule 3



(a)

Schedule 4

T_1	T_2
read(A) $A := A - 50$	read(A) $temp := A * 0.1$ $A := A - temp$ write(A) read(B)
write(A) read(B) $B := B + 50$ write(B) commit	$B := B + temp$ write(B) commit



Precedence graph for schedule 4.

The precedence graph for schedule 4 contains the edge $T_1 \rightarrow T_2$, because T_1 executes `read(A)` before T_2 executes `write(A)`.

It also contains the edge $T_2 \rightarrow T_1$, because T_2 executes `read(B)` before T_1 executes `write(B)`.

If the precedence graph for **S** has a cycle, then schedule **S** is not conflict serializable.

If the graph contains no cycles, then the **schedule S is conflict serializable**.

Thus, to test for conflict serializability, we need to construct the precedence graph and to invoke a cycle-detection algorithm. Cycle

- It is possible to have two schedules that produce the same outcome, but that are not conflict equivalent.
- For example, consider transaction T_5 , which transfers \$10 from account B to account A .
- Let us define schedule 8, this not conflict equivalent to the serial schedule $\langle T_1, T_5 \rangle$,
- since, in schedule 8, the $\text{write}(B)$ instruction of T_5 conflicts with the $\text{read}(B)$ instruction of T_1 .
- This creates an edge $T_5 \rightarrow T_1$ in the precedence graph. Similarly, we see that the $\text{write}(A)$ instruction of T_1 conflicts with the read instruction of T_5 creating an edge $T_1 \rightarrow T_5$.
- This shows that the precedence graph has a cycle and that schedule 8 is not serializable.
- However, the final values of accounts A and B after the execution of either schedule 8 or the serial schedule $\langle T_1, T_5 \rangle$ are the same—\$960 and \$2040, respectively. (Exercise)

Schedule 8.

T_1	T_5
$\text{read}(A)$ $A := A - 50$ $\text{write}(A)$	$\text{read}(B)$ $B := B - 10$ $\text{write}(B)$
$\text{read}(B)$ $B := B + 50$ $\text{write}(B)$	$\text{read}(A)$ $A := A + 10$ $\text{write}(A)$

Effect of transaction failures during concurrent execution.

- If a transaction T_i fails, for whatever reason, we need to undo the effect of this transaction to ensure the atomicity property of the transaction.
- In a system that allows concurrent execution, the atomicity property requires that any transaction T_j that is dependent on T_i (that is, T_j has read data written by T_i) is also aborted.
- To achieve this, we need to place restrictions on the type of schedules permitted in the system.
- what schedules are acceptable from the viewpoint of recovery from transaction failure.
 - **Recoverable Schedules**
 - **Cascadeless Schedules**

- **Recoverable Schedules**

- Consider the partial schedule 9, in which T_7 is a transaction that performs only one instruction: $\text{read}(A)$.
- We call this a **partial schedule** because we have not included a **commit** or **abort** operation for T_6 .
- Notice that T_7 commits immediately after executing the $\text{read}(A)$ instruction.
- Thus, T_7 commits while T_6 is still in the active state.
- Now suppose that T_6 fails before it commits. T_7 has read the value of data item A written by T_6 . Therefore, we say that T_7 is **dependent** on T_6 .
- Because of this, we must abort T_7 to ensure atomicity.
- However, T_7 has already committed and cannot be aborted.
- Thus, we have a situation where it is impossible to recover correctly from the failure of T_6 .

Schedule 9,

T_6	T_7
$\text{read}(A)$ $\text{write}(A)$	$\text{read}(A)$ commit
$\text{read}(B)$	

a nonrecoverable schedule.

A **recoverable schedule** is one where, for each pair of transactions T_i and T_j such that T_j reads a data item previously written by T_i , the commit operation of T_i appears before the commit operation of T_j .

For the example of schedule 9 to be recoverable, T_7 would have to delay committing until after T_6 commits.

- **Cascadeless Schedules**

- Even if a schedule is recoverable, to recover correctly from the failure of a transaction T_i , we may have to roll back several transactions.

- Such situations occur if many transactions have read data written by T_i .

- Illustration

T_8	T_9	T_{10}
read(A)		
read(B)		
write(A)		
	read(A)	
	write(A)	
		read(A)
abort		

Schedule 10.

- Transaction T_8 writes a value of A that is read by transaction T_9 . Transaction T_9 writes a value of A that is read by transaction T_{10} .
- Suppose that, at this point, T_8 fails. T_8 must be rolled back. Since T_9 is dependent on T_8 , T_9 must be rolled back. Since T_{10} is dependent on T_9 , T_{10} must be rolled back.
- This phenomenon, in which a single transaction failure leads to a series of transaction rollbacks, is called **cascading rollback**.

- Cascading rollback is undesirable, since it leads to the undoing of a significant amount of work.
- It is desirable to restrict the schedules to those where cascading rollbacks cannot occur. Such schedules are called *cascadeless* schedules.
- Formally, a **cascadeless schedule** is one where, for each pair of transactions T_i and T_j such that T_j reads a data item previously written by T_i , the commit operation of T_i appears before the read operation of T_j .
- It is easy to verify that every cascadeless schedule is also recoverable.

The End.//